

An Introduction to Programming in R

Drew Schmidt

November 13, 2015



About

This tutorial:

- This is not a “how to fit a linear model in R” tutorial.
- Slides and exercises are available at:
<http://wrathematics.info/handouts/bas2015.html>
- I try to stick to base R as much as possible, even when better alternatives exist.
- Mostly independent; we can skip something that's boring!

Me:

- I am a very productive R package developer.
- I mostly do not show you the things I make.
- My expertise is R in HPC (contact me if interested!)



1 Introduction

- What is R?
- Installing R
- Resources and Advice

2 R Basics

3 Programming

4 Closing

1 Introduction

- What is R?
- Installing R
- Resources and Advice


What is R?

- *lingua franca* for data analytics and statistical computing.
- Part programming language, part data analysis package.
- Dialect of S (Bell Labs).
- Syntax designed for data.

Who uses R?























Language Paradigms



Very Popular for a DSL!


IEEE Spectrum's 2014 Ranking of Programming Languages

| Language Rank | Types | Spectrum Ranking |
|---------------|---|------------------|
| 1. Java |    | 100.0 |
| 2. C |    | 99.3 |
| 3. C++ |    | 95.5 |
| 4. Python |   | 93.4 |
| 5. C# |    | 92.4 |
| 6. PHP |  | 84.7 |
| 7. Javascript |   | 84.4 |
| 8. Ruby |  | 78.8 |
| 9. R |  | 74.2 |
| 10. MATLAB |  | 72.9 |

See:

<http://spectrum.ieee.org/static/interactive-the-top-programming-languages#index>





At Build 2015 Microsoft CVP Joseph Sirosh called R the "language of data" and said *"if there is a single language that you choose to learn today .. let it be R"*.


1 Introduction

- What is R?
- Installing R
- Resources and Advice



Installing R

Go to <http://cran.r-project.org/> for a download (binary or source)



CRAN
[Mirrors](#)
[What's new?](#)
[Task Views](#)
[Search](#)

About R
[R Homepage](#)
[The R Journal](#)

Software
[R Sources](#)
[R Binaries](#)
[Packages](#)
[Other](#)

Documentation
[Manuals](#)
[FAQs](#)
[Contributed](#)

The Comprehensive R Archive Network

Download and Install R

Precompiled binary distributions of the base system and contributed packages, **Windows and Mac** users most likely want one of these versions of R:

- [Download R for Linux](#)
- [Download R for \(Mac\) OS X](#)
- [Download R for Windows](#)

R is part of many Linux distributions, you should check with your Linux package management system in addition to the link above.

Source Code for all Platforms

Windows and Mac users most likely want to download the precompiled binaries listed in the upper box, not the source code. The sources have to be compiled before you can use them. If you do not know what this means, you probably do not want to do it!

- The latest release (2013-05-16, Masked Marvel): [R-3.0.1.tar.gz](#), read [what's new](#) in the latest version.
- Sources of [R alpha and beta releases](#) (daily snapshots, created only in time periods before a planned release).
- Daily snapshots of current patched and development versions are

Installing R

Windows users should also install Rtools

<http://cran.r-project.org/bin/windows/Rtools/>

For a complete set of installation instructions, see

<http://www.r-pbd.org/install.html>

Installing Rstudio

Go to <http://www.rstudio.com/ide/>

The image shows the RStudio IDE website with the heading "Take control of your R code". Below the heading, it states: "RStudio is a free and open source integrated development environment for R. You can run it on your desktop (Windows, Mac, or Linux) or even over the web using RStudio Server." A "Download RStudio" button is visible, with a subtext "for Windows, Mac or Linux".

To the right of the text is a "Screencast" section titled "RStudio in 2 minutes" with a play button icon.

Below the text is a preview of the RStudio IDE interface. The interface shows a script editor with R code for loading the 'diamonds' dataset, summarizing it, and creating a faceted scatter plot of Price vs. Carat. The console shows the output of the summary and the plot. The plot is titled "Diamond Pricing" and shows a scatter plot of Price (y-axis, 0 to 10000) vs. Carat (x-axis, 0.000 to 0.000). The plot is faceted by clarity (VS1, VS2, VS3, VS4).

1 Introduction

- What is R?
- Installing R
- Resources and Advice

Important things we can't cover

- R and version control
- Developing R packages
- Performance/profiling
- Graphics/visualization

R Resources

- *The Art of R Programming* by Norm Matloff:
<http://nostarch.com/artofr.htm>
- *An Introduction to R* by Venables, Smith, and the R Core Team:
<http://cran.r-project.org/doc/manuals/R-intro.pdf>
- *The R Inferno* by Patrick Burns:
http://www.burns-stat.com/pages/Tutor/R_inferno.pdf
- Mathesaurus: <http://mathesaurus.sourceforge.net/>
- R programming for those coming from other languages: http://www.johndcook.com/R_language_for_programmers.html
- *aRrgh: a newcomer's (angry) guide to R*, by Tim Smith and Kevin Ushey: <http://tim-smith.us/arrgh/>

Tutorials

- *R Programming*, Coursera course through Johns Hopkins
<https://www.coursera.org/course/rprog>
- *Statistics One* Coursera course through Princeton
<https://www.coursera.org/course/stats1>
- *High Performance Computing with R*
<https://github.com/wrathematics/2015hpcRworkshop/blob/master/README.md>

Other Invaluable Resources

- *R Installation and Administration*:
<http://cran.r-project.org/doc/manuals/R-admin.html>
- *Task Views*: <http://cran.at.r-project.org/web/views>
- *Writing R Extensions*:
<http://cran.r-project.org/doc/manuals/R-exts.html>
- Mailing list archives: <http://tolstoy.newcastle.edu.au/R/>
- The [R] stackoverflow tag.
- The #rstats hashtag on Twitter.

Comments and Advice

- R is part statistics package, part programming language.
- R is slow; if you don't know what you're doing, it's *really* slow.
- There is an R help mailing list. Use stackoverflow instead....
- Learn to love the R help system.
- If something appears broken in core R, it's (probably) not them, it's you.
- Try to avoid “super functions”.

1 Introduction

2 R Basics

- R Basics
- All* About R Packages
- I/O
- Strings
- Dates
- Dealing with Dataframes

3 Programming

4 Closing



2 R Basics

- R Basics
- All* About R Packages
- I/O
- Strings
- Dates
- Dealing with Dataframes

Interacting with the R Terminal

The default thing for R to do is print/show:

```
1 1
2 1+1
3 print(1+1)
4 sum
5
6 + # ctrl+c to break
7 "+"
8 `+`
```

Arithmetic

```
1 7+4
2 7-4
3 7*4
4 7/4
5 7^4
6 7%4
```

Assignment

R naming rules can be quite lax. For all practical purposes:

- Start with a letter
- Should consist of letters (any case), numbers, .'s, and _'s
- Very strange things are possible, however...

```
1 m(list=ls())
```

```
1 "&*()" <- 3
2
3 "2" <- 1
4 `2` + 1
5 ## [1] 2
6
7 "\U0001f431" <- "even unicode"
8 cat(ls())
9 ## &*() 🐱 2
```


Assignment

```
1 # Local
2 myvar <- 1
3 myvar = 4
4
5 # Global
6 myvar <<- 2
7
8 # Whoever you want
9 assign(x="myvar", value=3)
```

Case and Spacing

R is case sensitive, but fairly lax about spacing:

```
1 x <- 1:5
2 x
3 X
4
5 1      + 1
6 sum(x)
7 sum ( x      )
8 s um(x)
```

Finding Help

- R has its own manual system.
- Most of the answers to your questions lie within.
- Find help using `?` or `help()`, or search across all help with `??`.

```
1 ?sum
2 ??sum
3 help("sum")
4
5 ?+
6 ?"+"
```

Ways to Run R

- Interactive: typing commands into the console.
- Batch: Rscript, R CMD BATCH.
- From an IDE/GUI.
- Batch from a GUI: `source()`

On the note of directories...

- Windows paths use `/` or `\`, e.g. `c:/myfile.r`
- Get current working directory: `getwd()`
- Set current working directory: `setwd("path/to/my/dir")`

2 R Basics

- R Basics
- All* About R Packages
- I/O
- Strings
- Dates
- Dealing with Dataframes

R Extensions

- R is great, but limited.
- Has a great package extension system.
- R doesn't do what you want? Make it!

"...if I don't know how to fix it, I can hire somebody else to fix it for me."
— Matt Dowle, developer of the `data.table` package.

Packages

- “R extensions”
- Comprehensive R Archive Network (CRAN).
- 7474 packages (`nrow(available.packages())`)
- CRAN Task Views <http://cran.r-project.org/web/views/>

Terminology

- A package is a collection of code usable by R.
- A library is a collection of packages.
- `library()` loads a package.
- Don't think too much about this...

Terminology

- We call the extension a *package*.
- *Packages* go into a *library*.
- **This is dumb and confusing, but there's nothing we can ever do about it.**

Example Confusion

- I wrote a library.
- I put the library in a package.
- I install the package ...into a library.
- I load the package with `library()` ???

BOOM



Installing R Packages

```
1 install.packages("devtools")
```

```
1 install.packages("devtools", lib="some/place/on/disk")
```

```
1 R CMD INSTALL devtools_1.6.tar.gz -l /some/place/on/disk
```

Repositories

- This basically assumes you're using CRAN.
- Lots of exciting development is happening outside of CRAN these days.
- Other binary package repositories: Bioconductor, R-forge (Windows)
- Other source repositories: GitHub, Bitbucket, ...

GitHub Binaries?



Karl Broman @kwbroman · Sep 27

.@hadleywickham @millerdl It would be great to have a service that would compile windows/mac binaries of #rstats pkgs on github.

← Reply ↻ Retweet ★ Favorite ... More



Hadley Wickham

@hadleywickham

+ Follow

@kwbroman @millerdl agreed. Plans are in motion

Installing Packages from Source

To install packages from source, you need some compilers:

- **Windows:** Install [Rtools](#).
- **Mac:** Install Xcode from the app store, possibly some other things from [here](#). If you need OpenMP, god help you.
- **Linux and FreeBSD:** You're good to go.

Installing R Packages

```
1 devtools::install_github("wrathematics/lineSampler")
```

Devtools Package Install Functions

| | | |
|-------------------|-------------------|-----------------|
| install_bitbucket | install_github | install_svn |
| install_deps | install_gitorious | install_url |
| install_git | install_local | install_version |

2 R Basics

- R Basics
- All* About R Packages
- I/O
- Strings
- Dates
- Dealing with Dataframes

I/O

| Input | Output |
|--------------------------|-----------------------------|
| <code>readcsv()</code> | <code>writecsv()</code> |
| <code>readtable()</code> | <code>writetable()</code> |
| <code>readBin()</code> | <code>writeBin()</code> |
| <code>readLines()</code> | <code>writelnLines()</code> |
| <code>save()</code> | <code>load()</code> |
| <code>saveRDS()</code> | <code>readRDS()</code> |
| <code>scan()</code> | |

“What about my really weird use-case?”

- NUMEROUS R packages
- XML, json, databases, ...
- See also: *R Data Import/Export* manual <https://cran.r-project.org/doc/manuals/r-release/R-data.html>

Reading and Writing a CSV

```
1 x <- matrix(1:30, 10)
2 write.csv(x, file="x.csv")
3 read.csv("x.csv")
4
5 write.csv(x, file="x.csv", row.names=FALSE)
6 read.csv("x.csv", header=TRUE)
```

The dreaded `stringsAsFactors`

- Default for R functions (e.g., `read.csv()`) is `stringsAsFactors=TRUE`.
- If you do not need to modify the “strings” (going straight to modeling), `stringsAsFactors=TRUE`.
- If you **DO** need to modify, `stringsAsFactors=FALSE`

Serializing

```
1 x <- letters
2 y <- 1:5
3 z <- list(list("a"), b=matrix(0))
4
5 save(x, y, z, file="objects.RData")
6 m(x)
7 m(y)
8 m(z)
9 x
10 load("objects.RData")
```

See also `saveRDS()`.

Why/Why Not Serialize?

Why:

- Serializing is a binary “as-is” format.
- Performance!

Why not:

- Only works with R (or something that can read R binary formats!)
- Endianness...

Some Notable Packages

General:

- rio <https://cran.r-project.org/web/packages/rio/index.html>

CSV and friends:

- data.table (fread()) <https://cran.r-project.org/web/packages/data.table/index.html>
- readr <https://cran.r-project.org/web/packages/readr/index.html>

Readers of proprietary things:

- haven (SAS, SPSS, and STATA) <https://cran.r-project.org/web/packages/haven/index.html>
- readxl (Excel) <https://cran.r-project.org/web/packages/readxl/index.html>

The `data()` Command

- Many packages (and R itself) bundle data.
- Load data with `data()`
- For a list of R's: `library(help="datasets")`

Example

```
1 head(iris)
2 head(mtcars)
3
4 data(package="ggplot2")
5 diamonds # error
6 data(diamonds)
7 head(diamonds)
```

2 R Basics

- R Basics
- All* About R Packages
- I/O
- **Strings**
- Dates
- Dealing with Dataframes

Strings

- “character” data (text).
- Internal storage scheme is complicated...
- Managing character data is its own course!
- Just the baby basics...

Quotes

- Three kinds: single quote ', double quote ", backtick `.
- Single and double create strings.
- Backticks access language objects (e.g., `+`).
- Escape within a string with a *single backslash*:

```
1 quoted_quote <- "\"\n"
2 quoted_quote
3 ## [1] "\"\n"
4 cat(quoted_quote)
5 ## "
```

Example 1

```
1 letters
2 ## [1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o"
   "p" "q" "r" "s"
3 ## [20] "t" "u" "v" "w" "x" "y" "z"
4
5 toupper(letters)
6 ## [1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O"
   "P" "Q" "R" "S"
7 ## [20] "T" "U" "V" "W" "X" "Y" "Z"
8
9 LETTERS
10 ## [1] "A" "B" "C" "D" "E" "F" "G" "H" "I" "J" "K" "L" "M" "N" "O"
   "P" "Q" "R" "S"
11 ## [20] "T" "U" "V" "W" "X" "Y" "Z"
12
13 tolower(LETTERS)
14 ## [1] "a" "b" "c" "d" "e" "f" "g" "h" "i" "j" "k" "l" "m" "n" "o"
   "p" "q" "r" "s"
15 ## [20] "t" "u" "v" "w" "x" "y" "z"
```

Example 2

```

1 x <- "Star Trek is objectively better than Star Wars"
2 strsplit(x, split=" ")
3 ## [[1]]
4 ## [1] "Star"          "Trek"          "is"           "objectively"
5      "better"
6
7 ## [6] "than"          "Star"          "Wars"
8
9 y <- unlist(strsplit(x, split=""))
10 y
11 ## [1] "S" "t" "a" "r" " " "T" "r" "e" "k" " " "i" "s" " " "o" "b"
12      "j" "e" "c" "t"
13 ## [20] "i" "v" "e" "l" "y" " " "b" "e" "t" "t" "e" "r" " " "t" "h"
14      "a" "n" " " "S"
15 ## [39] "t" "a" "r" " " "W" "a" "r" "s"
16
17 paste(rev(y), collapse="")
18 ## [1] "sraW ratS naht retteb ylevitcejbo si kerT ratS"

```


Example 3

```
1 paste(letters, letters)
2 ## [1] "a a" "b b" "c c" "d d" "e e" "f f" "g g" "h h" "i i" "j j"
   "k k" "l l"
3 ## [13] "m m" "n n" "o o" "p p" "q q" "r r" "s s" "t t" "u u" "v v"
   "w w" "x x"
4 ## [25] "y y" "z z"
5
6 paste(letters, letters, sep="")
7 ## [1] "aa" "bb" "cc" "dd" "ee" "ff" "gg" "hh" "ii" "jj" "kk" "ll"
   "mm" "nn" "oo"
8 ## [16] "pp" "qq" "rr" "ss" "tt" "uu" "vv" "ww" "xx" "yy" "zz"
9
10 paste(letters, letters, sep="", collapse="")
11 ## [1] "aabbccddeeffgghhiijjkkllmmnnnooppqqrrssttuuvvwwxxyyzz"
12
13 paste(paste(letters, collapse=""), paste(LETTERS, collapse=""),
14       sep="")
15 ## [1] "abcdefghijklmnopqrstuvwxyzABCDEFGHIJKLMNOPQRSTUVWXYZ"
```

Example 4

```
1 x <- rnorm(1000, mean=10)
2
3 paste("The mean of 'x' is:", mean(x))
4 ## [1] "The mean of 'x' is: 10.0157377724829"
5
6 cat(sprintf("mean:\t%.2f\nvar:\t%.2f\n", mean(x), var(x)))
7 ## mean:    10.02
8 ## var:     0.95
```

Regular Expressions

- `grep()`, `grepl()`
- `sub()`, `gsub()`
- `regexpr()`, `gregexpr()`
- Some others...

Some Notable Packages

- `stringi`
<https://cran.r-project.org/web/packages/stringi/index.html>
- `tm` <https://cran.r-project.org/web/packages/tm/index.html>

- 2 R Basics
 - R Basics
 - All* About R Packages
 - I/O
 - Strings
 - **Dates**
 - Dealing with Dataframes

Dates

- Dates managed in the UNIX style.
- Probably very alien for you...
- Just use lubridate.

Formatting Dates with Lubridate

```
1 rightnow <- Sys.time()
2 rightnow
3
4 library(lubridate)
5 year(rightnow)
6 ## [1] 2015
7
8 month(rightnow)
9 ## [1] 11
10 month(rightnow, label=TRUE)
11 ## [1] Nov
12
13 wday(rightnow, label=TRUE)
14 ## [1] Mon
```

And much more!

See the package vignette: <https://cran.r-project.org/web/packages/lubridate/vignettes/lubridate.html>

Example: Timezone Lookup

```
1 timezone <- function()  
2 {  
3   time <- Sys.time()  
4   ret <- list(timezone=format(time, format="%Z"),  
5               UTC.offset=format(time, format="%z"))  
6   class(ret) <- "tzlookup"  
7   return(ret)  
8 }  
9 print.tzlookup <- function(x)  
10 {  
11   maxlen <- max(sapply(names(x), nchar))  
12   spacenames <- c("timezone:", "UTC Offset:")  
13   cat(paste(spacenames, x, sep=" ", collapse="\n"), "\n")  
14 }  
15  
16 timezone()  
17 ## timezone:      EST  
18 ## UTC Offset: -0500
```

Some Notable Packages

- lubridate <https://cran.r-project.org/web/packages/lubridate/index.html>

- 2 R Basics
 - R Basics
 - All* About R Packages
 - I/O
 - Strings
 - Dates
 - Dealing with Dataframes

Motivation

- For many, dataframes are the most important/common object.
- ENORMOUS topic.
- We cover some important techniques.

Modifying Elements

```
1 ### Matrix-like notation
2 iris[1,1] <- 0
3 iris[, 1] <- 0
4
5 edit(iris)
```

Adding Variables

```
1 index <- 1:nrow(iris)
2 iris$index <- index
3 head(iris)
4
5 m(iris)
6 iris <- cbind(iris, index)
7 head(iris)
8
9 m(iris)
10 iris[, ncol(iris)+1] <- index
11 head(iris)
12
13 m(iris)
```

Adding Rows

```
1 nrow(iris)
2 set.seed(12345)
3 newrow <- iris[sample(nrow(iris), size=1), ]
4
5 iris <- rbind(iris, newrow)
6 nrow(iris)
7
8 iris[nrow(iris)+1, ] <- newrow
9 nrow(iris)
```

Adding Variables

```
1 index <- 1:nrow(iris)
2 iris$index <- index
3 head(iris)
4
5 m(iris)
6 head(iris)
7
8 iris <- cbind(iris, index)
9 head(iris)
10
11 m(iris)
```


Subsetting

```
1 subset(iris , Sepal.Length > 7)
2
3 subset(iris , Sepal.Length > 7 & Petal.Length < 6)
4
5 subset(iris , Species == "setosa" | Species == "versicolor")
```

Sorting/Ordering Rows

```
1 df1 <- iris[order(iris$Sepal.Length), ]  
2 head(df1)  
3  
4 df2 <- iris[order(-iris$Sepal.Length), ]  
5 head(df2)  
6  
7 df3 <- iris[order(-iris$Sepal.Length, iris$Sepal.Width), ]  
8 head(df3)
```

Sorting/Ordering Columns

```
1 iris <- iris[1:5, ]
2
3 iris[5:1]
4 set.seed(12345)
5 iris[sample(1:ncol(iris))]
```



```
6
7 iris[1]
8
9 iris[sort(colnames(iris))]
```



```
10
11 iris[sort(colnames(iris), decreasing=TRUE)]
```



```
12
13 m(iris)
```

Dealing with Duplicates

```
1 any(duplicated(iris))  
2 which(duplicated(iris))  
3  
4 nrow(unique(iris))
```

Applying Functions

```
1 colMeans(iris)
2 colMeans(iris[, -5])
3
4 apply(iris[, -5], MARGIN=2, FUN=median)
```

Some Notable Packages

Dataframe-like things:

- data.table <https://cran.r-project.org/web/packages/data.table/index.html>
- dplyr <https://cran.r-project.org/web/packages/dplyr/index.html>

Restructuring helpers:

- reshape2 <https://cran.r-project.org/web/packages/reshape2/index.html>
- tidyr <https://cran.r-project.org/web/packages/tidyr/index.html>
- broom <https://cran.r-project.org/web/packages/broom/index.html>

- 1 Introduction
- 2 R Basics
- 3 Programming
 - Type and Structure
 - Control Flow
 - Loops
 - Functions
 - Debugging
- 4 Closing

- 3 Programming
 - Type and Structure
 - Control Flow
 - Loops
 - Functions
 - Debugging

The R Language

- Storage: logical, int, double, double complex, character (strings)
- Structures: vector, matrix, array, list, dataframe, environment (hashtable)
- Caveats: (Logical) TRUE, FALSE, NA
- In fact, there's an NA for each type:
 - Integer: $-(2^{31} - 1)$
 - Double: value at address 0x7FF00000000007A2LL
- 3 official OOP systems, several unofficial ones.
- Several unofficial packages supporting C++.

Data Types

- logical
- int
- double
- double complex
- character

Data Types

R is *dynamically typed*. You do not have to declare what kind of data a variable is before you start using it:

```
1 x <- 1
2 typeof(x)
3 x <- 1:2
4 typeof(x)
```

Other

There are 4 “other” data “types”:

- Inf
- NaN
- NULL
- NA

Inf

Numerical infinity

```
1 Inf
2 ## [1] Inf
3
4 typeof(Inf)
5 ## [1] "double"
6
7 is.finite(Inf)
8 ## [1] FALSE
9
10 is.infinite(Inf)
11 ## [1] TRUE
12
13 Inf+Inf
14 ## [1] Inf
15
16 1/0
17 ## [1] Inf
```

NaN

Not a Number; numerical undefinedness.

```
1 NaN
2 ## [1] NaN
3
4 typeof(NaN)
5 ## [1] "double"
6
7 is.nan(NaN)
8 ## [1] TRUE
9
10 Inf - Inf
11 ## [1] NaN
12
13 sin(Inf)
14 ## [1] NaN
15 ## Warning message:
16 ## In sin(Inf) : NaNs produced
```

NULL

The null object; a sort of placeholder for something undefined. Like a non-numeric **NaN**.

```
1 NULL
2 ## NULL
3
4 typeof(NULL)
5 ## [1] "NULL"
6
7 is.null(NULL)
8 ## [1] TRUE
9
10 NULL+NULL
11 ## numeric(0)
```

NA

Missingness; not merely undefined, but *unknown*.

- Each type (logical, int, double, ...) has its own NA
- R is thus not boolean: TRUE, FALSE, NA
- Most R methods have ways of removing NA's

Data Structures

- vector
- matrix
- array
- factor
- list
- dataframe

Vectors

```
1 1:10
2 10:1
3
4 c(1, 3, 5, 7, 9)
5 seq(from=1, to=10, by=2)
6
7 x <- 1:5
8 length(x)
```

Matrices

```
1 matrix(1:10)
2 matrix(1:10, nrow=5)
3 matrix(1:10, ncol=5)
4
5 x <- 1:10
6 y <- as.matrix(x)
7 dim(x)
8 dim(y)
9 dim(x) <- c(1, 10)
10 x
```

Extraction

```
1 x <- matrix(1:30, 10)
2 x
3
4 x[-1, ]
5 x[, -1]
6 x[1:5, -1]
7 x[c(2,5,7), c(1,3)]
8
9 y <- x[, -2]
10 dim(y) <- NULL
11 y
```

Replacement

```
1 x <- matrix(1:30, 10)
2
3 x[1:5, ] <- 0
4 x[7, 3] <- NA
5 x
```

Factors

```
1 factor(1:5)
2 factor(c("a", "b", "b", "a", "c"))
3
4 x <- factor(-1:1)
5 x
6 as.numeric(x)
7 as.numeric(as.character(factor(-1:1)))
```

Dataframes

```
1 c(1, "a")
2 matrix(c(1, "a"))
3
4 x <- data.frame(1, "a")
5 x
6 x[1, 1]
7 is.numeric(x[1,1])
8 x[1, 2]
9 is.numeric(x[1,2])
10
11 data.frame(a=1:5,b=5:1)
```

Lists

- Super structures
- Items can be any structure (even other lists)
- Dataframe is really just a special list

Lists

```
1 list(1)
2 x <- list(list(1), "a")
3 x
4 x[[1]]
5
6 x <- list(a=list("b", 1), z=1:5)
7 x$z
```

Other Structures?

- stacks, heaps, queues, graphs, ...
- tl;dr: mostly no
- Hash table - environments
- Example deque <https://github.com/wrathematics/dequer>

3 Programming

- Type and Structure
- **Control Flow**
- Loops
- Functions
- Debugging

Logic

- Possible values are TRUE, FALSE, and NA
- Operations: &&, ||, &, |
- Comparators: ==, !=, <, <=, >, and >=

Logic

```
1 1==1
2 1!=1
3 1<1
4 1<=1
5
6 TRUE==FALSE
7
8 TRUE==1
9 FALSE==0
10
11 TRUE==T
12 FALSE==F
```

Logic

```
1 NA==NA
2 is.na(NA)
3
4 NULL==NULL
5 is.null(NULL)
6
7 NaN==NaN
8 is.nan(NaN)
9
10 Inf==Inf
11 is.infinite(Inf)
```

Logic

Vectors of logicals are evaluated element-wise:

```
1 x <- c(TRUE, FALSE, FALSE, TRUE)
2 x==TRUE
3 !x
```

Logic

Some very important functions for logical evaluations:

```
1 x <- c(TRUE, FALSE, FALSE, TRUE)
2 any(x)
3 all(x)
4 which(x)
5
6 y <- -5:5
7 which(y%%2==0)
8 y[which(y%%2==0)]
```


Comparisons

Comparing to...

- Comparing to NULL: `is.null(x)`
- Comparing to NA: `is.na(x)`
- Comparing to TRUE: `isTRUE(x)`

Comparing two...

- Tread carefully...: `x == y`
- Numerics: `all.equal(x, y)`
- EXACTLY THE SAME: `identical(x, y)`

Why all this trouble?

BECAUSE COMPUTERS ARE TERRIBLE

```
1 x <- 1
2 for (i in 1:10) x <- x - .1
3 x
```

Another famous example

```
1 sprintf("%.17f", 0.1+0.2)
2 ## [1] "0.30000000000000004"
```

This even has its own website! <http://0.30000000000000004.com/>

3 Programming

- Type and Structure
- Control Flow
- **Loops**
- Functions
- Debugging

Loops

- Instruction set to be repeatedly executed until some condition is satisfied (possibly forever).
- R has `for()` and `while()`.
- `for()`: Iterates over a list or vector
- `while()`: Performs operations until logical condition is satisfied.

for Loop

```
1 for (i in 1:10){
2   cat(i, " ")
3 }
4 ## 1 2 3 4 5 6 7 8 9 10
5
6 x <- matrix(1:30, 10)
7 colmax <- numeric(3)
8 for (i in 1:3){
9   colmax[i] <- max(x[, i])
10 }
11
12 colmax
13 ## [1] 10 20 30
```

while Loop

```
1 i <- 1
2 while (i < 11){
3   print(i)
4   i <- i+1
5 }
6
7 n <- 2
8 i <- 1
9 while (n < 1000){
10   n <- n^2
11   i <- i*2
12 }
13
14 n
15 ## [1] 65536
16 i
17 ## [1] 16
```

The *ply Family

- `apply()`: Apply function across “margin” (dimension) of matrix.
- `lapply()`: Apply function to input data object; returns a list.
- `sapply()`: Same as `sapply()` but
- `mapply()`: Multivariate `sapply()`.
- `vapply()`: Essentially `sapply()`, sometimes faster.
- `tapply()`: Applying a function to a subset of a vector.
- ...

We will only discuss the first 3.

What is it?

- Syntax to loop without writing a loop.
- Inspired by functional programming.
- **Not the same thing as vectorization** (but it looks vectorized).

apply

```
1 x <- matrix(1:30, 10)
2 x
3
4 apply(X=x, MARGIN=1, FUN=min)
5 ## [1] 1 2 3 4 5 6 7 8 9 10
6 apply(X=x, MARGIN=2, FUN=min)
7 ## [1] 1 11 21
8
9 out <- numeric(3)
10 for (i in 1:3){
11   out[i] <- min(x[, i])
12 }
13 out
14 ## [1] 1 11 21
```

lapply and sapply

```
1 x <- 1:5
2 lapply(X=x, FUN=sqrt)
3 ## [[1]]
4 ## [1] 1
5 ##
6 ## [[2]]
7 ## [1] 1.414214
8 ##
9 ## [[3]]
10 ## [1] 1.732051
11 ##
12 ## [[4]]
13 ## [1] 2
14 ##
15 ## [[5]]
16 ## [1] 2.236068
17
18 sapply(X=x, FUN=sqrt)
19 ## [1] 1.000000 1.414214 1.732051 2.000000 2.236068
```

*ply Functions Internally

- `apply()`: A `for()` loop.
- `lapply()`: Internal R voodoo; faster than a loop.
- `sapply()`: Essentially the same as `lapply()`.

When does all this choice matter?

- loops are slow.
- `apply()` is sugar for an R for loop.
- `lapply()` different, often faster.
- Vectorization is fastest of all.

Loop Speeds

```
1 x <- 100000:1
2
3 system.time({ # No initialization
4   sin <- numeric(0)
5   for (i in 1:length(x)){
6     sin[i] <- sin(x[i])
7   }
8   sin
9 })
10 ##      user  system elapsed
11 ## 17.320   1.072   18.376
12
13 system.time({ # With initialization
14   sin <- numeric(length(x))
15   for (i in 1:length(x)){
16     sin[i] <- sin(x[i])
17   }
18   sin
19 })
20 ##      user  system elapsed
21 ##  0.172   0.000   0.171
```

Loop Speeds

```
1 system.time(lapply(x, sin))
2 ## user system elapsed
3 ## 0.056 0.004 0.059
4
5 system.time(sapply(x, sin))
6 ## user system elapsed
7 ## 0.048 0.004 0.053
8
9 system.time(sin(x))
10 ## user system elapsed
11 ## 0.004 0.000 0.004
```

3 Programming

- Type and Structure
- Control Flow
- Loops
- **Functions**
- Debugging

Functions

- Self-contained input/output machine.
- Reusable blocks of code.
- First class objects.
- All functions have returns!
- *Evaluating the Design of the R language*,
<http://r.cs.purdue.edu/pub/ecoop12.pdf>

Some Basic Rules

- Can not modify data in place* (multi-return with a list)
- Arguments can have defaults, specified with `=`.
- By default, no printing occurs (have to say `print(x)`).
- The last “thing done” is the return.
- People often use `NULL` or `invisible(NULL)` if return is unimportant.

Functions: Example 1

```
1 f <- function(x)
2 {
3   ret <- x+1
4   return(ret)
5 }
6
7 f(1)
8 ## [1] 2
9 f(2)
10 ## [1] 3
11 f(5)
12 ## [1] 6
13 f
14 ## function(x)
15 ## {
16 ##   ret <- x+1
17 ##   return(ret)
18 ## }
```

Functions: Example 2

```
1 f <- function (a, b)
2 {
3   a - b
4 }
5
6 f(a=1, b=2)
7 f(1, 2)
8 f(b=1, a=2)
9 f(b=1, 2)
10 f(1)
11 f(matrix(1:4, nrow=2), matrix(4:1, nrow=2))
```

Functions: Example 3

For complicated returns (especially of mixed type/class), use a list:

```
1 g <- function (a, b)
2 {
3   plus <- a+b
4   minus <- a-b
5   return(list(plus, minus))
6 }
7
8 g(5, 2)
9 g(1, 0)
10 g(f(2, 6), 2)
```

Functions: Example 4

R allows parameter defaults in functions

```
1 h <- function (a=1, b=2)
2 {
3   return(b-a)
4 }
5
6 h
7 h()
8 h(2, 1)
```

Recursive Functions

A *recursive function* is one that calls itself:

```
1 f <- function(n)
2 {
3   if (n==1)
4     return(1)
5   else {
6     if (n%%2==0)
7       return(f(n/2))
8     else
9       return(f(3*n+1))
10  }
11 }
12
13 f(22)
14 ## [1] 1
15 f(237)
16 ## [1] 1
17 sapply(1:37, f)
18 ##      [1] 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1 1
      [1] 1 1 1 1 1 1 1
```

3 Programming

- Type and Structure
- Control Flow
- Loops
- Functions
- Debugging

Debugging R Code

- Very broad topic ...
- We'll hit the highlights.
- For more examples, see:
cran.r-project.org/doc/manuals/R-exts.html#Debugging
- Debugging compiled code called by R (valgrind, gdb, ...) also possible...

Object Inspection Tools

- `print()`
- `str()`
- `unclass()`

Object Inspection Tools: print()

Basic printing:

```
1 x <- matrix(1:10, nrow=2)
2 print(x)
3 ##      [,1] [,2] [,3] [,4] [,5]
4 ## [1,]    1    3    5    7    9
5 ## [2,]    2    4    6    8   10
6
7 x
8 ##      [,1] [,2] [,3] [,4] [,5]
9 ## [1,]    1    3    5    7    9
10 ## [2,]    2    4    6    8   10
```

Object Inspection Tools: str()

Examining the structure of an R object:

```
1 x <- matrix(1:10, nrow=2)
2
3 str(x)
4 ##  int [1:2, 1:5] 1 2 3 4 5 6 7 8 9 10
```

Object Inspection Tools: unclass()

Exposing all data with unclass():

```
1 df <- data.frame(x=rnorm(10), y=rnorm(10))
2 mdl <- lm(y~x, data=df) ### That's a "tilde" character
3
4 mdl
5 print(mdl)
6
7 str(mdl)
8
9 unclass(mdl)
```

The R Debugger

- `debug()`
- `debugonce()`
- `undebug()`

Using The R Debugger

- 1 Declare function to be debugged: `debug(foo)`
- 2 Call function: `foo(arg1, arg2, ...)`
 - `next:` Enter or n followed by Enter.
 - `break:` Halt execution and exit debugging: Q.
 - `exit:` Continue execution and exit debugging: C.
- 3 Call `undebug()` to stop debugging

Using the Debugger

Example Debugger Interaction

```
> f <- function(x){y <- z+1;z <- y*2;z}
> f(1)
Error in f(1) : object 'z' not found
> debug(f)
> f(1)
debugging in: f(1)
debug at #1: {
  y <- z + 1
  z <- y * 2
  z
}
Browse[2]>
debug at #1: y <- z + 1
Browse[2]>
Error in f(1) : object 'z' not found
>
```


- 1 Introduction
- 2 R Basics
- 3 Programming
- 4 Closing

Thanks so much for attending!

Questions?

