

A Beginner's Guide to Programming in R

Drew Schmidt

June 9, 2015



About This Tutorial

Slides and Exercises

The slides and exercises are available at:

<http://wrathematics.info/handouts/dance2015.html>



Following Along

- Log in to `arronax.nics.utk.edu` (you may wish to forward X11)
- Run `module load r/3.0.1`
- Run R



1 Introduction to R

- What is R?
- Resources and Advice

2 Basics

3 Programming

4 Project

1 Introduction to R

- What is R?
- Resources and Advice

What is R?

- *lingua franca* for data analytics and statistical computing.
- Part programming language, part data analysis package.
- Dialect of S (Bell Labs).
- Syntax designed for data.

Who uses R?

ORACLE®

MERCK

KICKSTARTER

Bank of America



eBay

The
New York
Times

NIST



Pfizer

mozilla
FOUNDATION

FDA

ORBITZ

cfpb



NOVARTIS

LLOYD'S



bing


Google



Zillow®



Language Paradigms























The R Language

- R is slow; if you don't know what you're doing, it's *really* slow.
- High-level scripting language.
- Syntax designed for data: models are first-class constructs, missingness is built into the core of the language, ...

But you can't deny its popularity!


IEEE Spectrum's 2014 Ranking of Programming Languages

Language Rank	Types	Spectrum Ranking
1. Java	  	100.0
2. C	  	99.3
3. C++	  	95.5
4. Python	 	93.4
5. C#	  	92.4
6. PHP		84.7
7. Javascript	 	84.4
8. Ruby		78.8
9. R		74.2
10. MATLAB		72.9

See:

<http://spectrum.ieee.org/static/interactive-the-top-programming-languages#index>





At Build 2015 Microsoft CVP Joseph Sirosh called R the "language of data" and said *"if there is a single language that you choose to learn today .. let it be R"*.

1 Introduction to R

- What is R?
- Resources and Advice

R Resources

- *The Art of R Programming* by Norm Matloff:
<http://nostarch.com/artofr.htm>
- *An Introduction to R* by Venables, Smith, and the R Core Team:
<http://cran.r-project.org/doc/manuals/R-intro.pdf>
- *The R Inferno* by Patrick Burns:
http://www.burns-stat.com/pages/Tutor/R_inferno.pdf
- Mathesaurus: <http://mathesaurus.sourceforge.net/>
- R programming for those coming from other languages: http://www.johndcook.com/R_language_for_programmers.html
- *aRrgh: a newcomer's (angry) guide to R*, by Tim Smith and Kevin Ushey: <http://tim-smith.us/arrgh/>

Tutorials

- *R Programming*, Coursera course through Johns Hopkins
<https://www.coursera.org/course/rprog>
- *Statistics One* Coursera course through Princeton
<https://www.coursera.org/course/stats1>
- *High Performance Computing with R* <https://github.com/wrathematics/2015hpcRworkshop/blob/master/README.md>

Other Invaluable Resources

- *R Installation and Administration*:
<http://cran.r-project.org/doc/manuals/R-admin.html>
- *Task Views*: <http://cran.at.r-project.org/web/views>
- *Writing R Extensions*:
<http://cran.r-project.org/doc/manuals/R-exts.html>
- Mailing list archives: <http://tolstoy.newcastle.edu.au/R/>
- The [R] stackoverflow tag.
- The #rstats hashtag on Twitter.

Comments and Advice

- R is part statistics package, part programming language.
- There are always 100 ways to do anything, only one of them efficient.
- R is slow; if you don't know what you're doing, it's *really* slow.
- There is an R help mailing list. Use stackoverflow instead. . . .
- Learn to love the R help system.
- If something appears broken in core R, it's probably not a bug (it's you).
- Just because something is on the CRAN does not mean it's of any value (or even functional).
- Indent your code.

Comments and Advice

- Be consistent.
- Generally try to use instructive names for things.
- Make your code concise, but not obtuse (don't play golf).
- Try to avoid “super functions”. Breaking up complicated ideas into modular pieces can help with readability, debugging, reusability, etc.
- Google has an R style guide: <https://google-styleguide.googlecode.com/svn/trunk/google-r-style.html>
- The R community has few computer scientists.

1 Introduction to R

2 Basics

- R Basics
- Strings
- I/O
- Plotting
- Packages

3 Programming

4 Project

2 Basics

- R Basics
- Strings
- I/O
- Plotting
- Packages

Starting R

- 1 Interactive: R
- 2 Batch: Rscript

Interacting with the Terminal

The default method in R is `show()`, which usually amounts to printing:

```
1 1
2 1+1
3 print(1+1)
4 sum
5
6 + # ctrl+c to break
7 "+"
8 '+'
```

Arithmetic

```
1 7+4
2 7-4
3 7*4
4 7/4
5 7^4
6 7%%4
```

Assignment

R naming rules can be quite lax. For all practical purposes:

- Start with a letter
- Should consist of letters (any case), numbers, .'s, and _'s
- Very strange things are possible, however...

```
1 "&*()" <- 3
2 "&*()"
3 "2" <- 1
4 '2' + 1
5 " _ " <- "even unicode"
```

Assignment

```
1 myvar <- 1
2 myvar
3
4 myvar <<- 2
5 myvar
6
7 assign(x="myvar", value=3)
8 myvar
9
10 myvar = 4
11 myvar
```


Case and Spacing

R is case sensitive, but fairly lax about spacing:

```
1 x <- 1:5
2 x
3 X
4
5 1      +    1
6 sum(x)
7 sum ( x    )
8 s um(x)
```

Finding Help

- R has its own manual system.
- Most of the answers to your questions lie within.
- Find help using `?` or `help()`, or search across all help with `??`.

```
1 ?sum
2 ??sum
3 help("sum")
4
5 ?+ # ctrl+c to break
6 ?"+"
```

RNG

R contains many powerful random number generators:

Beta	Binomial	Cauchy
Chi-square	Exponential	F
Gamma	Geometric	Hypergeometric
Logistic	Log Normal	Negative Binomial
Normal	Poisson	Student t
Uniform	Weibull	Wilcoxon Rank Sum

```
1 runif(5, min=0, max=10)
2 rnorm(5, mean=0, sd=1)
3 rgamma(5, shape=1)
4
5 set.seed(10)
6 runif(5)
7 set.seed(10)
8 runif(5)
```

2 Basics

- R Basics
- **Strings**
- I/O
- Plotting
- Packages

Strings

- “Character” data (text).
- Internal storage scheme is complicated. . .
- Mostly works like most high-level languages.
- Can have a vector, matrix, dataframe, or list of strings.

Strings: Example 1

```
1 letters
2 toupper(letters)
3 LETTERS
4 tolower(LETTERS)
```

Strings: Example 2

```
1 x <- "Star Trek is objectively better than Star Wars"
2 strsplit(x, split=" ")
3
4 y <- unlist(strsplit(x, split=" "))
5 y
6 paste(rev(y), collapse=" ")
```

Strings: Example 3

```
1 paste(letters, letters)
2 paste(letters, letters, sep=" ")
3 paste(letters, letters, sep=" ", collapse=" ")
4
5 paste(paste(letters, collapse=" "), paste(LETTERS, collapse=" "),
      sep=" ")
```


Strings: Example 4

```
1 x <- rnorm(1000, mean=10)
2
3 paste("The mean of 'x' is:", mean(x))
4
5 cat(sprintf("mean:\t%.2f\nvar:\t%.2f\n", mean(x), var(x)))
```

Regular Expressions

- `grep()`, `grep1()`
- `sub()`, `gsub()`
- Some others...

2 Basics

- R Basics
- Strings
- I/O
- Plotting
- Packages

I/O

- `write.csv()`, `read.csv()`,
- `write.table()`, `read.table()`
- `scan()`
- `save()` and `load()`

Reading and Writing a CSV

```
1 x <- matrix(1:30, 10)
2 write.csv(x, file="x.csv")
3 read.csv("x.csv")
4
5 write.csv(x, file="x.csv", row.names=F)
6 read.csv("x.csv", header=T)
```

Serializing

```
1 x <- letters
2 y <- 1:5
3 z <- list(list("a"), b=matrix(0))
4
5 save(x, y, z, file="robjects.RData")
6 rm(x)
7 rm(y)
8 rm(z)
9 x
10 load("robjects.RData")
```

See also `saveRDS()`.

2 Basics

- R Basics
- Strings
- I/O
- **Plotting**
- Packages

Plotting

- R is *world class* for graphics and visualization.
- Biggest package for plotting is **ggplot2** by Hadley Wickham.
- We will focus on some examples from core R. . .

Scatterplots

```
1 x <- 1:10
2 y <- rnorm(10)
3 plot(x, y)
4 text(x=5.5, y=0, "MY PLOT", col='red')
```

Piecharts

```
1 pie(1:5)
```

Barplots

```
1 barplot(1:5)
```

Histograms

```
1 hist(rnorm(500))
```

Saving Plots

```
1 pdf("myscatterplot.pdf")
2 plot(rnorm(4), 1:4)
3 dev.off()
4
5 png("mybarplot.png")
6 barplot(1:10)
7 dev.off()
```

Plotting

- These are all fairly “cookie-cutter”.
- These can do more complicated things, but you’re only making things difficult for yourself. . .
- Best to learn a full graphics package.
- The biggest are **ggplot2** and **lattice** (I recommend ggplot. . .)

Lattice: <http://www.statmethods.net/advgraphs/trellis.html>

ggplot2: <http://www.statmethods.net/advgraphs/ggplot2.html>

2 Basics

- R Basics
- Strings
- I/O
- Plotting
- Packages

Packages

- Comprehensive R Archive Network (CRAN).
- 6735 packages.
- CRAN Task Views <http://cran.r-project.org/web/views/>

Installing Packages

From CRAN

```
1 install.packages("devtools")
```

From GitHub

```
1 library("devtools")  
2 install_github("wrathematics/lineSampler")
```

Namespaces

All packages have namespaces:

```
1 devtools::install_github("wrathematics/lineSampler")
```

Libraries

- A package is a collection of code usable by R.
- A library is a collection of packages.
- `library()` loads a package.
- Don't think too much about this...

1 Introduction to R

2 Basics

3 Programming

- Type and Structures

- Type
- Structures

- Control Flow

- Logic
- Loops
- *ply

- Functions

- Debugging

- Debugging R Code

- The R Debugger

- Debugging Compiled Code Called by R Code



3 Programming

- Type and Structures
 - Type
 - Structures
- Control Flow
 - Logic
 - Loops
 - *ply
- Functions
- Debugging
 - Debugging R Code
- The R Debugger
 - Debugging Compiled Code Called by R Code

The R Language

- Storage: logical, int, double, double complex, character (strings)
- Structures: vector, matrix, array, list, dataframe
- Caveats: (Logical) TRUE, FALSE, NA
- In fact, there's an NA for each type:
 - Integer: $-(2^{-31} - 1)$
 - Double: value at address 0x7FF00000000007A2LL
- 3 official OOP systems (none of them work like anything you're used to), several unofficial ones.
- No official support for C++; several unofficial packages supporting C++.

Data Types

- logical
- int
- double
- double complex
- character

Data Types

R is *dynamically typed*. You do not have to declare what kind of data a variable is before you start using it:

```
1 x <- 1
2 typeof(x)
3 x <- 1:2
4 typeof(x)
```


Other

There are 4 “other” data “types”:

- Inf
- NaN
- NULL
- NA

Inf

Numerical infinity

```
1 Inf
2 typeof(Inf)
3 is.finite(Inf)
4 is.infinite(Inf)
5 Inf+Inf
6 1/0
```

NaN

Not a Number; numerical undefinedness.

```
1 NaN
2 typeof(NaN)
3 is.nan(NaN)
4 Inf - Inf
5 sin(Inf)
```

NULL

The null object; a sort of placeholder for something undefined. Like a non-numeric NaN.

```
1 NULL
2 typeof(NULL)
3 is.null(NULL)
4 NULL+NULL
```

NA

Missingness; not merely undefined, but *unknown*.

- Each type (logical, int, double, ...) has its own NA
- R is thus not boolean: TRUE, FALSE, NA
- Most R methods have ways of removing NA's

Data Structures

- vector
- matrix
- array
- factor
- list
- dataframe

Vectors

```
1 1:10
2 10:1
3
4 c(1, 3, 5, 7, 9)
5 seq(from=1, to=10, by=2)
6
7 x <- 1:5
8 length(x)
9 is.vector(x)
```

Matrices

```
1 matrix(1:10)
2 matrix(1:10, nrow=5)
3 matrix(1:10, ncol=5)
4
5 x <- 1:10
6 y <- as.matrix(x)
7 dim(x)
8 dim(y)
9 dim(x) <- c(1, 10)
10 x
```


Extraction

```
1 x <- matrix(1:30, 10)
2 x
3
4 x[-1, ]
5 x[, -1]
6 x[1:5, -1]
7 x[c(2,5,7), c(1,3)]
8
9 y <- x[, -2]
10 dim(y) <- NULL
11 y
```

Replacement

```
1 x <- matrix(1:30, 10)
2 x
3
4 x[1:5, ] <- 0
5 x[7, 3] <- NA
6 x
```

Factors

```
1 factor(1:5)
2 factor(c("a", "b", "b", "a", "c"))
3
4 x <- factor(-1:1)
5 x
6 as.numeric(x)
7 as.numeric(as.character(factor(-1:1)))
```

Dataframes

```
1 c(1, "a")
2 matrix(c(1, "a"))
3
4 x <- data.frame(1, "a")
5 x
6 x[1, 1]
7 is.numeric(x[1,1])
8 x[1, 2]
9 is.numeric(x[1,2])
10
11 data.frame(a=1:5, b=5:1)
```

Lists

- Super structures
- Items can be any structure (even other lists)
- Dataframe is really just a special list

Lists

```
1 list(1)
2 x <- list(list(1), "a")
3 x
4 x[[1]]
5
6 x <- list(a=list("b", 1), z=1:5)
7 x$z
```

Other Structures?

- stacks, heaps, queues, graphs, ...
- tl;dr: no
- Example deque <https://github.com/wrathematics/dequer>

3 Programming

- Type and Structures
 - Type
 - Structures
- Control Flow
 - Logic
 - Loops
 - *ply
- Functions
- Debugging
 - Debugging R Code
- The R Debugger
 - Debugging Compiled Code Called by R Code

Logic

- Possible values are TRUE, FALSE, and NA
- Relational operations: ==, !=, <, <=, >, and >=

Logic

```
1 1==1
2 1!=1
3 1<1
4 1<=1
5
6 TRUE==FALSE
7
8 TRUE==1
9 FALSE==0
10
11 TRUE==T
12 FALSE==F
```

Logic

```
1 NA==NA
2 is.na(NA)
3
4 NULL==NULL
5 is.null(NULL)
6
7 NaN==NaN
8 is.nan(NaN)
9
10 Inf==Inf
11 is.infinite(Inf)
```

Logic

Vectors of logicals are evaluated element-wise:

```
1 x <- c(T, F, F, T, T, F)
2 x==T
3 !x
```

Logic

Some very important functions for logical evaluations:

```
1 x <- c(T, F, F, T, T, F)
2 any(x)
3 all(x)
4 which(x)
5
6 y <- -5:5
7 which(y%%2==0)
8 y[which(y%%2==0)]
```

Loops

- Instruction set to be repeatedly executed until some condition is satisfied (possibly forever).
- R has `for()` and `while()`.
- `for()`: Iterates over a list or vector
- `while()`: Performs operations until logical condition is satisfied.

for Loop

```
1 for (i in 1:10){  
2   print(i)  
3 }  
4  
5 x <- matrix(1:30, 10)  
6 colmax <- numeric(3)  
7 for (i in 1:3){  
8   colmax[i] <- max(x[, i])  
9 }  
10 colmax
```

while Loop

```
1 i <- 1
2 while (i < 11){
3   print(i)
4   i <- i+1
5 }
6
7 n <- 2
8 i <- 1
9 while (n < 1000){
10   n <- n^2
11   i <- i*2
12 }
13 n
14 i
```


The *ply Family

- `apply()`: Apply function across “margin” (dimension) of matrix.
- `lapply()`: Apply function to input data object; returns a list.
- `sapply()`: Same as `sapply()` but
- `mapply()`: Multivariate `sapply()`.
- `vapply()`: Essentially `sapply()`, sometimes faster.
- `tapply()`: Applying a function to a subset of a vector.

We will only discuss the first 3.

apply

```
1 x <- matrix(1:30, 10)
2 x
3 apply(X=x, MARGIN=1, FUN=min)
4 apply(X=x, MARGIN=2, FUN=min)
5
6 out <- numeric(3)
7 for (i in 1:3){
8   out[i] <- min(x[, i])
9 }
10 out
```

lapply and sapply

```
1 x <- 1:5
2 lapply(X=x, FUN=sqrt)
3
4 sapply(X=x, FUN=sqrt)
```

*ply Functions Internally

- `apply()`: A `for()` loop.
- `lapply()`: Internal R voodoo; faster than a loop.
- `sapply()`: Essentially the same as `lapply()`.

When does all this choice matter?

- loops are slow.
- `apply()` is sugar for an R for loop.
- `lapply()` different, often faster.
- Vectorization is fastest of all.

Loop Speeds

```
1 # No initialization
2 system.time({
3   x <- 5000:1
4   sin <- numeric(0)
5   for (i in 1:length(x)){
6     sin[i] <- sin(x[i])
7   }
8   sin
9 })
10
11 # With initialization
12 system.time({
13   x <- 5000:1
14   sin <- numeric(length(x))
15   for (i in 1:length(x)){
16     sin[i] <- sin(x[i])
17   }
18   sin
19 })
```



Loop Speeds

```
1 # lapply
2 system.time(lapply(x, sin))
3 system.time(sapply(x, sin))
4
5 # vectorized
6 system.time(sin(x))
```

3 Programming

- Type and Structures
 - Type
 - Structures
- Control Flow
 - Logic
 - Loops
 - *ply
- **Functions**
- Debugging
 - Debugging R Code
- The R Debugger
 - Debugging Compiled Code Called by R Code

Functions

- Self-contained input/output machine.
- Reusable blocks of code.
- First class objects.
- Can not modify in place*; use a list.
- *Evaluating the Design of the R language*,
<http://r.cs.purdue.edu/pub/ecoop12.pdf>

Functions: Example 1

```
1 f <- function(x)
2 {
3   ret <- x+1
4   return(ret)
5 }
6
7 f(1)
8 f(2)
9 f(5)
10 f
```

Functions: Example 2

```
1 f <- function (a, b)
2 {
3   a - b
4 }
5
6 f(a =1, b=2)
7 f(1, 2)
8 f(b=1, a=2)
9 f(b=1, 2)
10 f(1)
11 f(matrix(1:4, ncol=2), matrix(4:1, nrow=2))
```

Functions: Example 3

For complicated returns (especially of mixed type/class), use a list:

```
1 g <- function (a, b)
2 {
3   return(list(plus=a+b, minus=a-b, times=a*b, division=a/b))
4 }
5
6 g(5, 2)
7 g(1, 0)
8 g(f(2, 6), 2)
```

Functions: Example 4

R allows parameter defaults in functions

```
1 h <- function (a=1, b=2)
2 {
3   return(b-a)
4 }
5
6 h
7 h()
8 h(2, 1)
```

Recursive Functions

A *recursive function* is one that calls itself:

```
1 f <- function(n)
2 {
3   if (n==1)
4     return(1)
5   else {
6     if (n%%2==0)
7       return(f(n/2))
8     else
9       return(f(3*n+1))
10  }
11 }
```

3 Programming

- Type and Structures
 - Type
 - Structures
- Control Flow
 - Logic
 - Loops
 - *ply
- Functions
- Debugging
 - Debugging R Code
- The R Debugger
 - Debugging Compiled Code Called by R Code

Debugging R Code

- Very broad topic ...
- We'll hit the highlights.
- For more examples, see:
cran.r-project.org/doc/manuals/R-exts.html#Debugging

Object Inspection Tools

- `print()`
- `str()`
- `unclass()`

Object Inspection Tools: print()

Basic printing:

```
1 > x <- matrix(1:10, nrow=2)
2 > print(x)
3      [,1] [,2] [,3] [,4] [,5]
4 [1,]    1    3    5    7    9
5 [2,]    2    4    6    8   10
6 > x
7      [,1] [,2] [,3] [,4] [,5]
8 [1,]    1    3    5    7    9
9 [2,]    2    4    6    8   10
```

Object Inspection Tools: `str()`

Examining the structure of an R object:

```
1 > x <- matrix(1:10, nrow=2)
2 > str(x)
3 int [1:2, 1:5] 1 2 3 4 5 6 7 8 9 10
```

Object Inspection Tools: unclass()

Exposing all data with unclass():

```
1 df <- data.frame(x=rnorm(10), y=rnorm(10))
2 mdl <- lm(y~x, data=df) ### That's a "tilde" character
3
4 mdl
5 print(mdl)
6
7 str(mdl)
8
9 unclass(mdl)
```

3 Programming

- Type and Structures
 - Type
 - Structures
- Control Flow
 - Logic
 - Loops
 - *ply
- Functions
- Debugging
 - Debugging R Code
- The R Debugger
 - Debugging Compiled Code Called by R Code

The R Debugger

- `debug()`
- `debugonce()`
- `undebug()`

Using The R Debugger

- 1 Declare function to be debugged: `debug(foo)`
- 2 Call function: `foo(arg1, arg2, ...)`
 - `next:` Enter or n followed by Enter.
 - `break:` Halt execution and exit debugging: Q.
 - `exit:` Continue execution and exit debugging: c.
- 3 Call `undebug()` to stop debugging


Using the Debugger

Example Debugger Interaction

```
1 > f <- function(x){y <- z+1;z <- y*2;z}
2 > f(1)
3 Error in f(1) : object 'z' not found
4 > debug(f)
5 > f(1)
6 debugging in: f(1)
7 debug at #1: {
8     y <- z + 1
9     z <- y * 2
10    z
11 }
12 Browse[2]>
13 debug at #1: y <- z + 1
14 Browse[2]>
15 Error in f(1) : object 'z' not found
16 >
```


Debugging Compiled Code

- Reasonably easy to use gdb and Valgrind (from command line).
- gdb — The GNU Debugger; general purpose debugging.
- Valgrind — Memory debugger.
- For gdb, start R interactively.
- For Valgrind, need a batch script.



Debugging with gdb

Suppose we have:

- R function: `fooR()`
- Calls the C function: `fooC()`

We can debug `fooC()` via `gdb` by executing the following from a shell:

```
1 R -d gdb
2 b fooC
3 signal 0
4 fooR(10)
```

Debugging with Valgrind

Put the R code you wish to profile in `myscript.r` and execute the following from a shell:

```
1 R -d "valgrind --tool=memcheck --leak-check=full" --vanilla <  
    myscript.r
```

- 1 Introduction to R
- 2 Basics
- 3 Programming
- 4 Project
 - Getting Started
 - Using xts

4 Project

- Getting Started
- Using xts

Getting Started

- Log in to `arronax.nics.utk.edu`
- Run `module load r/3.0.1`
- Run `cd /lustre/medusa/proj/dance/data/day_0129_csv`

Managing the Data

- `ls`
- `ls | grep Unit989`

Managing the Data

reconfigure.sh

```
1 #!/bin/sh
2
3 for file in `ls | grep Unit989`; do
4     cat $file >> Unit989.csv
5 done
```

```
1 $ chmod +x reconfigure.sh
2 $ ./reconfigure.sh
3 $ wc -l Unit989.csv
```


Managing the Data

reconfigure.sh

```
1 #!/bin/sh
2
3 for folder in `ls -d */ | grep -E
4     "day_[0-9][0-9][0-9][0-9].csv"`; do
5     for file in `ls $folder | grep Unit989`; do
6         cat $folder/$file >> Unit989.csv
7     done
8 done
```

```
1 $ chmod +x reconfigure.sh
2 $ ./reconfigure.sh
3 $ wc -l Unit989.csv
```

Loading

```
1 file <- "Unit989-20150129-000102.csv"
2 file.info(file)
3 data <- read.csv(file)
4 head(data)
5 str(data)
6 data <- read.csv(file, header=FALSE, stringsAsFactors=FALSE)
```

Inspecting

```
1 head(data)
2 tail(data)
3 dim(data)
4 str(data)
5 summary(data)
6 sd(data$V5)
7 object.size(data)
```

4 Project

- Getting Started
- Using xts

xts: The eXtensible Time Series package

```
1 library(xts)
```

Time Series

```
1 date <- as.POSIXlt(paste(data$V1, data$V2))
2 x <- xts(data$V4, date)
3 head(x)
4 str(x)
```

Time Series

```
1 first(x, "3 seconds")
2 first(x, "1 minute")
3 last(first(x, "1 minute"), "3 seconds")
```

Time Series

```
1 ndays(x)
2 nhours(x)
3 nminutes(x)
```


Time Series

```
1 to.daily(x)
2 to.hourly(x)
3 to.minutes30(x)
4 to.minutes15(x)
```

Time Series

```
1 endpoints(x, "hours")  
2 period.apply(x, INDEX=endpoints(x, "hours"), FUN=mean)
```

Time Series

```
1 min30 <- to.minutes30(x)
2
3 plot(x)
4 plot(min30)
5
6 plot(as.zoo(min30))
7 plot(as.zoo, plot.type="single")
8 plot(as.zoo, plot.type="single", col=1:4)
```